

# CHURN ANALYSIS

Predictive Modelling

Baijayanti  
Chakraborty  
PGP-BABI

1.Project Objective	2
2.Exploratory Data Analysis	2
2.1 Install and load the needed libraries.	2
2.2 Set the working directory and check the basic statistics of the dataset.	10
2.3 The descriptive analytics of the dataset.	12
2.3.1. Univariate and Bivariate analysis	15
2.3.2.Data Visualizations	16
2.3.3. Insights gained from the Exploratory Data Analysis can be concluded as :-	25
3.Model Creation	26
3.1.Logistic Regression	26
3.2.K-Nearest Neighbour(KNN) :-	52
3.3.Naive Bayes	53
4.Model Performance	56

## 1. Project Objective

Customer Churn is a burning problem for Telecom companies. In this project, we simulate one such case of customer churn where we work on a data of postpaid customers with a contract. The data has information about the customer usage behavior, contract details and the payment details. The data also indicates which were the customers who canceled their service. Based on this past data, we need to build a model which can predict whether a customer will cancel their service in the future or not.

One is expected to do the following :

### 1. EDA

- How does the data looks like, Univariate and bivariate analysis. Plots and charts which illustrate the relationships between variables
- Look out for outliers and missing values
- Check for multicollinearity & treat it
- Summarize the insights you get from EDA

### 2. Build Models and compare them to get to the best one

- Logistic Regression
- KNN
- Naive Bayes
- Model Comparison using Model Performance metrics & Interpretation

### 3. Actionable Insights

- Interpretation & Recommendations from the best model

## 2. Exploratory Data Analysis

### 2.1 Install and load the needed libraries.

```
#clean the global environment
rm(list = ls())
#set the working directory
#setwd("~/Desktop/PGP-BABI/Predictive Modelling/week3-frequency based")

#Load the Libraries
library(readxl)

## Warning: package 'readxl' was built under R version 3.6.1

library(DataExplorer)

## Warning: package 'DataExplorer' was built under R version 3.6.1

library(caret)

## Warning: package 'caret' was built under R version 3.6.1

## Loading required package: lattice
```

```
## Loading required package: ggplot2
## Warning: package 'ggplot2' was built under R version 3.6.1
library(psych)
## Warning: package 'psych' was built under R version 3.6.1
##
## Attaching package: 'psych'
## The following objects are masked from 'package:ggplot2':
##
##      %+%, alpha
library(GGally)
## Warning: package 'GGally' was built under R version 3.6.1
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
library(corrplot)
## Warning: package 'corrplot' was built under R version 3.6.1
## corrplot 0.84 loaded
library(earth)
## Warning: package 'earth' was built under R version 3.6.1
## Loading required package: Formula
## Loading required package: plotmo
## Warning: package 'plotmo' was built under R version 3.6.1
## Loading required package: plotrix
##
## Attaching package: 'plotrix'
## The following object is masked from 'package:psych':
##
##      rescale
## Loading required package: TeachingDemos
library(varImp)
## Warning: package 'varImp' was built under R version 3.6.1
## Loading required package: measures
## Warning: package 'measures' was built under R version 3.6.1
```

```
##
## Attaching package: 'measures'

## The following object is masked from 'package:psych':
##
##      AUC

## The following objects are masked from 'package:caret':
##
##      MAE, RMSE

## Loading required package: party
## Warning: package 'party' was built under R version 3.6.1
## Loading required package: grid
## Loading required package: mvtnorm
## Loading required package: modeltools
## Loading required package: stats4
## Loading required package: strucchange
## Warning: package 'strucchange' was built under R version 3.6.1
## Loading required package: zoo
## Warning: package 'zoo' was built under R version 3.6.1

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric

## Loading required package: sandwich
## Warning: package 'sandwich' was built under R version 3.6.1

##
## Attaching package: 'varImp'

## The following object is masked from 'package:caret':
##
##      varImp

library(IMTest)

## Warning: package 'IMTest' was built under R version 3.6.1
## Loading required package: ltm
## Warning: package 'ltm' was built under R version 3.6.1
```

```
## Loading required package: MASS
## Warning: package 'MASS' was built under R version 3.6.1
## Loading required package: msm
## Warning: package 'msm' was built under R version 3.6.1
## Loading required package: polycor
## Warning: package 'polycor' was built under R version 3.6.1
##
## Attaching package: 'polycor'
## The following object is masked from 'package:psych':
##
##     polyserial
##
## Attaching package: 'ltm'
## The following object is masked from 'package:psych':
##
##     factor.scores
library(pscl)
## Warning: package 'pscl' was built under R version 3.6.1
## Classes and Methods for R developed in the
## Political Science Computational Laboratory
## Department of Political Science
## Stanford University
## Simon Jackman
## hurdle and zeroinfl functions by Achim Zeileis
library(e1071)
## Warning: package 'e1071' was built under R version 3.6.1
library(caret)
library(pROC)
## Warning: package 'pROC' was built under R version 3.6.1
## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
library(class)
```

```
## Warning: package 'class' was built under R version 3.6.1
library(gmodels)
## Warning: package 'gmodels' was built under R version 3.6.1
##
## Attaching package: 'gmodels'
## The following object is masked from 'package:PROC':
##
##      ci
library(car)
## Warning: package 'car' was built under R version 3.6.1
## Loading required package: carData
##
## Attaching package: 'car'
## The following object is masked from 'package:modeltools':
##
##      Predict
## The following object is masked from 'package:psych':
##
##      logit
library(ROCR)
## Warning: package 'ROCR' was built under R version 3.6.1
## Loading required package: gplots
## Warning: package 'gplots' was built under R version 3.6.1
##
## Attaching package: 'gplots'
## The following object is masked from 'package:plotrix':
##
##      plotCI
## The following object is masked from 'package:stats':
##
##      lowess
library(blorr)
## Warning: package 'blorr' was built under R version 3.6.1
library(class)
library(car)
library(caret)
```

```
library(class)
library(devtools)

## Warning: package 'devtools' was built under R version 3.6.1
## Loading required package: usethis
## Warning: package 'usethis' was built under R version 3.6.1

library(e1071)
library(ggplot2)
library(Hmisc)

## Warning: package 'Hmisc' was built under R version 3.6.1
## Loading required package: survival

##
## Attaching package: 'survival'

## The following object is masked from 'package:caret':
##
##   cluster

##
## Attaching package: 'Hmisc'

## The following object is masked from 'package:e1071':
##
##   impute

## The following objects are masked from 'package:TeachingDemos':
##
##   cnvrt.coords, subplot

## The following object is masked from 'package:psych':
##
##   describe

## The following objects are masked from 'package:base':
##
##   format.pval, units

library(klaR)

## Warning: package 'klaR' was built under R version 3.6.1

## This version of Shiny is designed to work with 'htmlwidgets'
## >= 1.5.
##   Please upgrade via install.packages('htmlwidgets').

##
## Attaching package: 'klaR'
```



```
## The following object is masked from 'package:TeachingDemos':  
##  
##      triplot  
  
library(klaR)  
library(MASS)  
library(plyr)  
  
## Warning: package 'plyr' was built under R version 3.6.1  
  
##  
## Attaching package: 'plyr'  
  
## The following objects are masked from 'package:Hmisc':  
##  
##      is.discrete, summarize  
  
## The following object is masked from 'package:modeltools':  
##  
##      empty  
  
library(scatterplot3d)  
library(SDMTools)  
  
## Warning: package 'SDMTools' was built under R version 3.6.1  
  
## Registered S3 method overwritten by 'R.oo':  
##      method          from  
##      throw.default R.methodsS3  
  
##  
## Attaching package: 'SDMTools'  
  
## The following object is masked from 'package:PROC':  
##  
##      auc  
  
## The following objects are masked from 'package:caret':  
##  
##      sensitivity, specificity  
  
library(dplyr)  
  
## Warning: package 'dplyr' was built under R version 3.6.1  
  
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:plyr':  
##  
##      arrange, count, desc, failwith, id, mutate, rename,  
##      summarise,  
##      summarize
```

```
## The following objects are masked from 'package:Hmisc':
##
##   src, summarize

## The following object is masked from 'package:car':
##
##   recode

## The following object is masked from 'package:MASS':
##
##   select

## The following object is masked from 'package:GGally':
##
##   nasa

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ElemStatLearn)

## Warning: package 'ElemStatLearn' was built under R version
3.6.1

##
## Attaching package: 'ElemStatLearn'

## The following object is masked from 'package:plyr':
##
##   ozone

library(caTools)

## Warning: package 'caTools' was built under R version 3.6.1

library(boot)

## Warning: package 'boot' was built under R version 3.6.1

##
## Attaching package: 'boot'

## The following object is masked from 'package:survival':
##
##   aml

## The following object is masked from 'package:car':
##
##   logit
```

```
## The following object is masked from 'package:msm':
##
##     cav

## The following object is masked from 'package:psych':
##
##     logit

## The following object is masked from 'package:lattice':
##
##     melanoma
```

## 2.2 Set the working directory and check the basic statistics of the dataset.

```
#read the dataset
cell = read_excel("Cellphone.xlsx", sheet = 2)
#view the data set
head(cell, 7)

## # A tibble: 7 x 11
##   Churn AccountWeeks ContractRenewal DataPlan DataUsage
##   <dbl>         <dbl>          <dbl>    <dbl>    <dbl>
## 1     0          128              1      1      2.7
## 2     0          107              1      1      3.7
## 3     0          137              1      0      0
## 4     0           84              0      0      0
## 5     0           75              0      0      0
## 6     0          118              0      0      0
## 7     0          121              1      1      2.03

## # ... with 5 more variables: DayMins <dbl>, DayCalls <dbl>,
## #   MonthlyCharge <dbl>, OverageFee <dbl>, RoamMins <dbl>

#changing the column names to increase the readability
colnames(cell)

## [1] "Churn"           "AccountWeeks"    "ContractRenewal"
## [4] "DataPlan"        "DataUsage"       "CustServCalls"
## [7] "DayMins"         "DayCalls"        "MonthlyCharge"
## [10] "OverageFee"      "RoamMins"

names(cell)[2] = "AccWeeks"
names(cell)[3] = "ContRenew"
names(cell)[4] = "Plan"
```

```

names(cell)[5] = "Usage"
names(cell)[6] = "CScalls"
names(cell)[7] = "Min/Day"
names(cell)[8] = "Call/Day"
names(cell)[9] = "Charge/Month"
names(cell)[10] = "Over.Fee"
names(cell)[11] = "RoamingMins"
#Basic data summary, Univariate, Bivariate analysis, graphs
summary(cell)

##      Churn      AccWeeks      ContRenew      Plan
## Min.   :0.0000  Min.    : 1.0  Min.   :0.0000  Min.
##      :0.0000
## 1st Qu.:0.0000  1st Qu.: 74.0  1st Qu.:1.0000  1st
##      Qu.:0.0000
## Median :0.0000  Median :101.0  Median :1.0000  Median
##      :0.0000
## Mean   :0.1449  Mean   :101.1  Mean   :0.9031  Mean
##      :0.2766
## 3rd Qu.:0.0000  3rd Qu.:127.0  3rd Qu.:1.0000  3rd
##      Qu.:1.0000
## Max.   :1.0000  Max.    :243.0  Max.    :1.0000  Max.
##      :1.0000
##      Usage      CScalls      Min/Day      Call/Day
## Min.   :0.0000  Min.    :0.000  Min.    : 0.0  Min.   :
##      0.0
## 1st Qu.:0.0000  1st Qu.:1.000  1st Qu.:143.7  1st Qu.:
##      87.0
## Median :0.0000  Median :1.000  Median :179.4  Median
##      :101.0
## Mean   :0.8165  Mean   :1.563  Mean   :179.8  Mean
##      :100.4
## 3rd Qu.:1.7800  3rd Qu.:2.000  3rd Qu.:216.4  3rd
##      Qu.:114.0
## Max.   :5.4000  Max.    :9.000  Max.    :350.8  Max.
##      :165.0
##      Charge/Month      Over.Fee      RoamingMins
## Min.   : 14.00  Min.    : 0.00  Min.    : 0.00
## 1st Qu.: 45.00  1st Qu.: 8.33  1st Qu.: 8.50
## Median : 53.50  Median :10.07  Median :10.30
## Mean   : 56.31  Mean   :10.05  Mean   :10.24
## 3rd Qu.: 66.20  3rd Qu.:11.77  3rd Qu.:12.10
## Max.   :111.30  Max.    :18.19  Max.    :20.00

str(cell)

## Classes 'tbl_df', 'tbl' and 'data.frame': 3333 obs. of 11
##      variables:
## $ Churn      : num  0 0 0 0 0 0 0 0 0 0 ...
## $ AccWeeks   : num  128 107 137 84 75 118 121 147 117 141
##      ...
## $ ContRenew  : num  1 1 1 0 0 0 1 0 1 0 ...

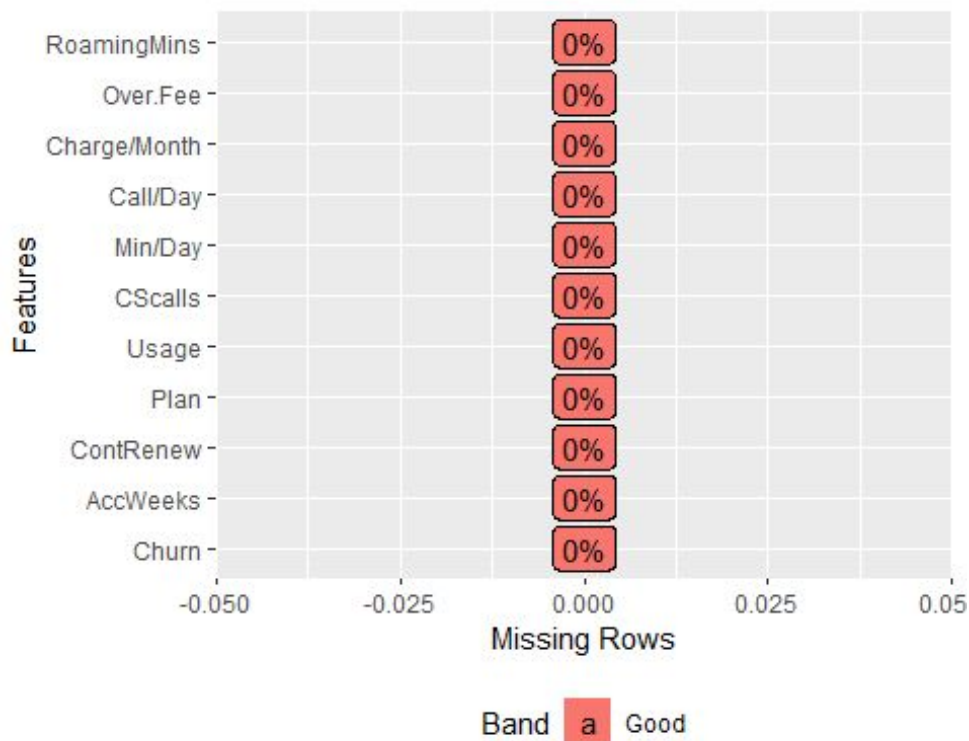
```

```
## $ Plan      : num  1 1 0 0 0 0 1 0 0 1 ...
## $ Usage     : num  2.7 3.7 0 0 0 0 2.03 0 0.19 3.02 ...
## $ CScalls   : num  1 1 0 2 3 0 3 0 1 0 ...
## $ Min/Day   : num  265 162 243 299 167 ...
## $ Call/Day  : num  110 123 114 71 113 98 88 79 97 84 ...
## $ Charge/Month: num  89 82 52 57 41 57 87.3 36 63.9 93.2 ...
## $ Over.Fee  : num  9.87 9.78 6.06 3.1 7.42 ...
## $ RoamingMins : num  10 13.7 12.2 6.6 10.1 6.3 7.5 7.1 8.7
11.2 ...

sum(is.na(cell)) #No null values in the whole dataset

## [1] 0

#graph for null value cross check
plot_missing(cell)
```



## 2.3 The descriptive analytics of the dataset.

*#descriptive analytics*

```
describe(cell)
```

```
## cell
##
## 11 Variables      3333 Observations
##
-----
-
## Churn
##      n missing distinct    Info      Sum      Mean      Gmd
## 3333      0         2    0.372    483    0.1449    0.2479
```

```

##
##
-----
-
## AccWeeks
##      n missing distinct      Info      Mean      Gmd      .05      .10
##    3333      0      212      1    101.1    45.01      35      50
##      .25      .50      .75      .90      .95
##      74      101      127      152      167
##
## lowest :   1   2   3   4   5, highest: 221 224 225 232 243
##
-----
-
## ContRenew
##      n missing distinct      Info      Sum      Mean      Gmd
##    3333      0      2    0.263    3010    0.9031    0.1751
##
##
-----
-
## Plan
##      n missing distinct      Info      Sum      Mean      Gmd
##    3333      0      2      0.6     922    0.2766    0.4003
##
##
-----
-
## Usage
##      n missing distinct      Info      Mean      Gmd      .05      .10
##    3333      0      174    0.839    0.8165    1.202      0.00      0.00
##      .25      .50      .75      .90      .95
##    0.00      0.00      1.78      3.05      3.46
##
## lowest : 0.00 0.11 0.12 0.13 0.14, highest: 4.59 4.64 4.73 4.75 5.40
##
-----
-
## CScalls
##      n missing distinct      Info      Mean      Gmd      .05      .10
##    3333      0      10    0.932    1.563    1.392      0      0
##      .25      .50      .75      .90      .95
##      1      1      2      3      4
##
## Value      0      1      2      3      4      5      6      7      8      9
## Frequency  697  1181  759  429  166  66  22  9  2  2
## Proportion 0.209 0.354 0.228 0.129 0.050 0.020 0.007 0.003 0.001 0.001
##
-----
-
## Min/Day

```

```

##          n missing distinct      Info      Mean      Gmd      .05      .10
##      3333         0      1667         1      179.8      61.46      89.92     110.32
##          .25      .50      .75      .90      .95
##      143.70     179.40     216.40     249.58     270.74
##
## lowest :    0.0    2.6    7.8    7.9  12.5, highest: 335.5 337.4 345.3
346.8 350.8
##
-----
-
## Call/Day
##          n missing distinct      Info      Mean      Gmd      .05      .10
##      3333         0       119         1     100.4      22.59      67.0      74.2
##          .25      .50      .75      .90      .95
##      87.0      101.0     114.0     126.0     133.0
##
## lowest :    0  30  35  36  40, highest: 157 158 160 163 165
##
-----
-
## Charge/Month
##          n missing distinct      Info      Mean      Gmd      .05      .10
##      3333         0       656         1     56.31     18.35     33.26     38.00
##          .25      .50      .75      .90      .95
##      45.00     53.50     66.20     80.50     87.80
##
## lowest :  14.0  15.7  16.0  17.0  19.0, highest: 108.3 108.6 108.7
110.0 111.3
##
-----
-
## Over.Fee
##          n missing distinct      Info      Mean      Gmd      .05      .10
##      3333         0     1024         1     10.05      2.86      5.94      6.84
##          .25      .50      .75      .90      .95
##      8.33     10.07     11.77     13.29     14.22
##
## lowest :  0.00  1.56  2.11  2.13  2.20, highest: 17.55 17.58 17.71
18.09 18.19
##
-----
-
## RoamingMins
##          n missing distinct      Info      Mean      Gmd      .05      .10
##      3333         0       162         1     10.24      3.114      5.7      6.7
##          .25      .50      .75      .90      .95
##      8.5      10.3      12.1      13.7      14.7
##
## lowest :  0.0  1.1  1.3  2.0  2.1, highest: 18.2 18.3 18.4 18.9 20.0
##

```

### 2.3.1. Univariate and Bivariate analysis

1. **Univariate data** –This type of data consists of only one variable. The analysis of univariate data is thus the simplest form of analysis since the information deals with only one quantity that changes. It does not deal with causes or relationships and the main purpose of the analysis is to describe the data and find patterns that exist within it
2. **Bivariate data** –This type of data involves two different variables. The analysis of this type of data deals with causes and relationships and the analysis is done to find out the relationship among the two variables.

Keeping this definition we see that our dataset is a Bivariate one since we see that more than one variable is involved in getting the relationship for the customer churn analysis. The “Churn” factor depends on other variables like “Plan”, “Usage”, “Cost of the renewal” etc.

*#Correlation Matrix*

```
corrplot(corr = cor(cell), method = "number" , type = "upper")
```



```
mat = cor(cell)
print(mat)
```

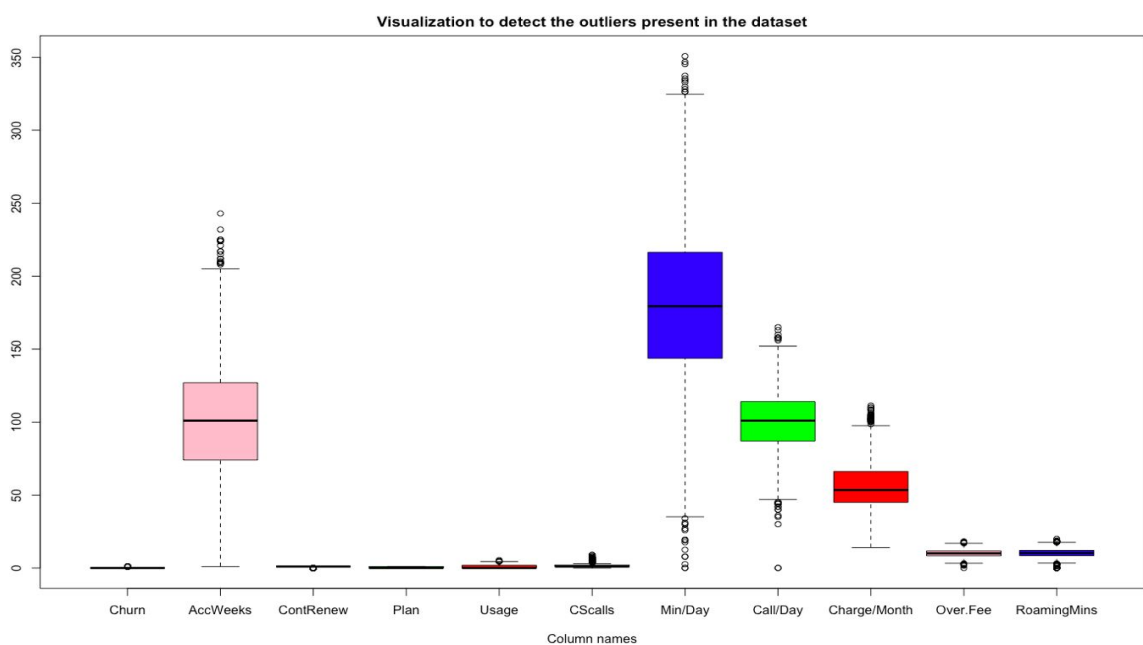
```
##           Churn      AccWeeks      ContRenew      Plan
## Churn      1.00000000  0.016540742 -0.259851847 -0.102148141
## AccWeeks   0.01654074  1.000000000 -0.024734655  0.002918409
## ContRenew  -0.25985185 -0.024734655  1.000000000 -0.006006371
## Plan       -0.10214814  0.002918409 -0.006006371  1.000000000
## Usage      -0.08719451  0.014390757 -0.019222913  0.945981734
## CScalls     0.20875000 -0.003795939  0.024521956 -0.017823944
## Min/Day     0.20515083  0.006216021 -0.049395824 -0.001684069
## Call/Day    0.01845931  0.038469882 -0.003754626 -0.011085902
## Charge/Month 0.07231271  0.012580670 -0.047291399  0.737489653
## Over.Fee    0.09281243 -0.006749462 -0.019104644  0.021525559
## RoamingMins 0.06823878  0.009513902 -0.045870743 -0.001317871
##           Usage      CScalls      Min/Day      Call/Day
```



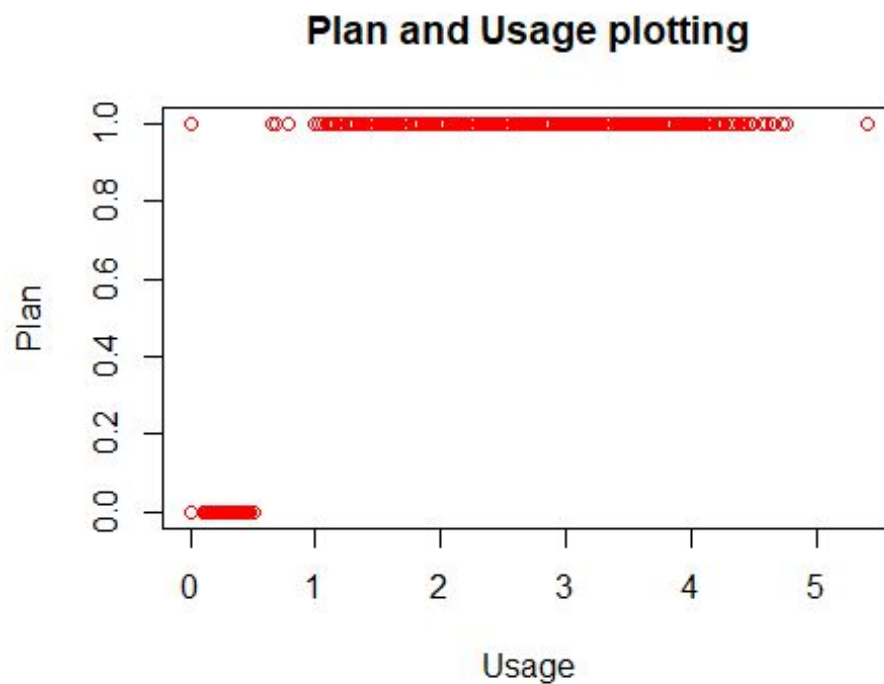
```
## Churn -0.087194509 0.208749999 0.205150829 0.018459312
## AccWeeks 0.014390757 -0.003795939 0.006216021 0.038469882
## ContRenew -0.019222913 0.024521956 -0.049395824 -0.003754626
## Plan 0.945981734 -0.017823944 -0.001684069 -0.011085902
## Usage 1.000000000 -0.021722518 0.003175951 -0.007962079
## CScalls -0.021722518 1.000000000 -0.013423186 -0.018941930
## Min/Day 0.003175951 -0.013423186 1.000000000 0.006750414
## Call/Day -0.007962079 -0.018941930 0.006750414 1.000000000
## Charge/Month 0.781660429 -0.028016853 0.567967924 -0.007963218
## Over.Fee 0.019637372 -0.012964219 0.007038214 -0.021448602
## RoamingMins 0.162745576 -0.009639680 -0.010154586 0.021564794
## Charge/Month Over.Fee RoamingMins
## Churn 0.072312711 0.092812426 0.068238776
## AccWeeks 0.012580670 -0.006749462 0.009513902
## ContRenew -0.047291399 -0.019104644 -0.045870743
## Plan 0.737489653 0.021525559 -0.001317871
## Usage 0.781660429 0.019637372 0.162745576
## CScalls -0.028016853 -0.012964219 -0.009639680
## Min/Day 0.567967924 0.007038214 -0.010154586
## Call/Day -0.007963218 -0.021448602 0.021564794
## Charge/Month 1.000000000 0.281766048 0.117432607
## Over.Fee 0.281766048 1.000000000 -0.011023336
## RoamingMins 0.117432607 -0.011023336 1.000000000
```

### 2.3.2. Data Visualizations


```
boxplot(cell,main = "Visualization to detect the outliers present in
the dataset",xlab = "Column names",col =
c("red","pink","blue","green"))
```




```
plot(cell$Usage,cell$Plan,main = "Plan and Usage plotting" , xlab =  
"Usage" , ylab = "Plan" , col = "red")
```



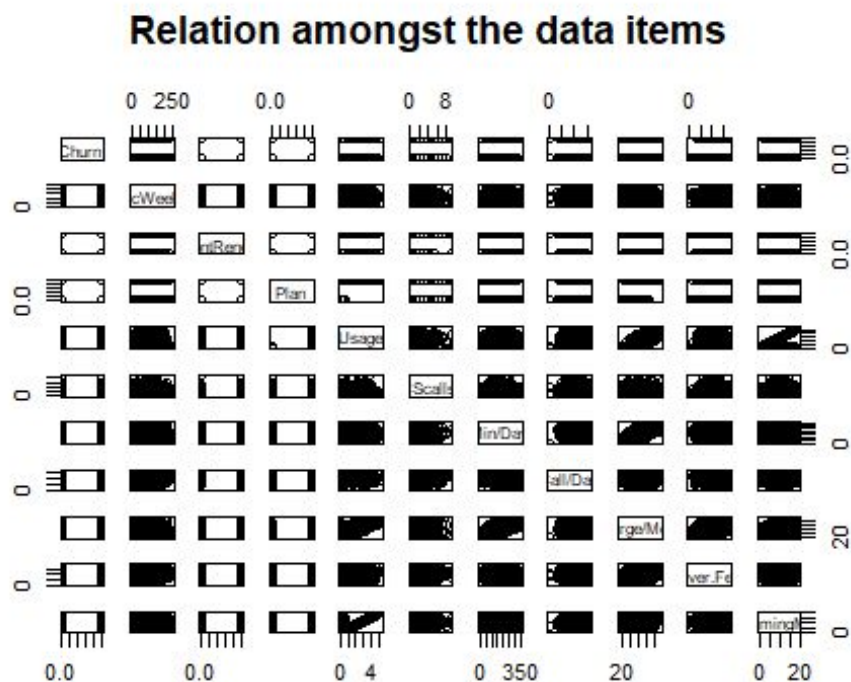
```
plot(cell$`Charge/Month`,cell$Usage , main = "Monthly Charge and Usage  
plan" , xlab = "Monthly Charge",ylab = "Data Usage",col = "blue")
```



```
plot(cell$`Charge/Month`,cell$Plan , main = "Monthly Charge and Data Plan"
, xlab = "Monthly Charge",ylab = "Data Plan",col = "green")
```



```
pairs(cell , main = "Relation amongst the data items")
```

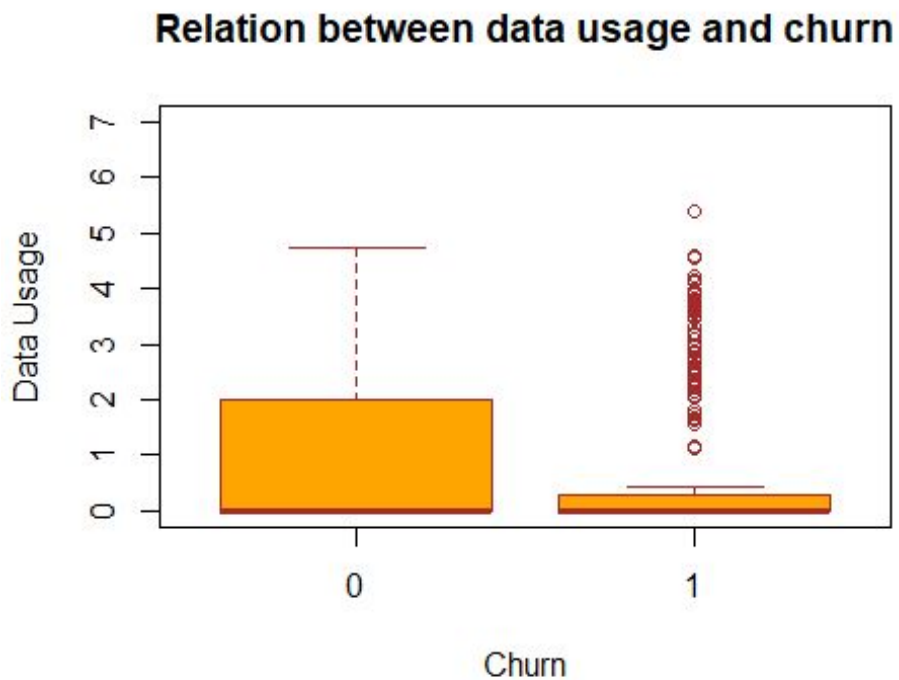


```
boxplot(Usage~Churn,
data = cell,
```

```

main = "Relation between data usage and churn",
xlab = "Churn",
ylab = "Data Usage",
col = "orange",
border = "brown",
ylim = c(0, 7)
)

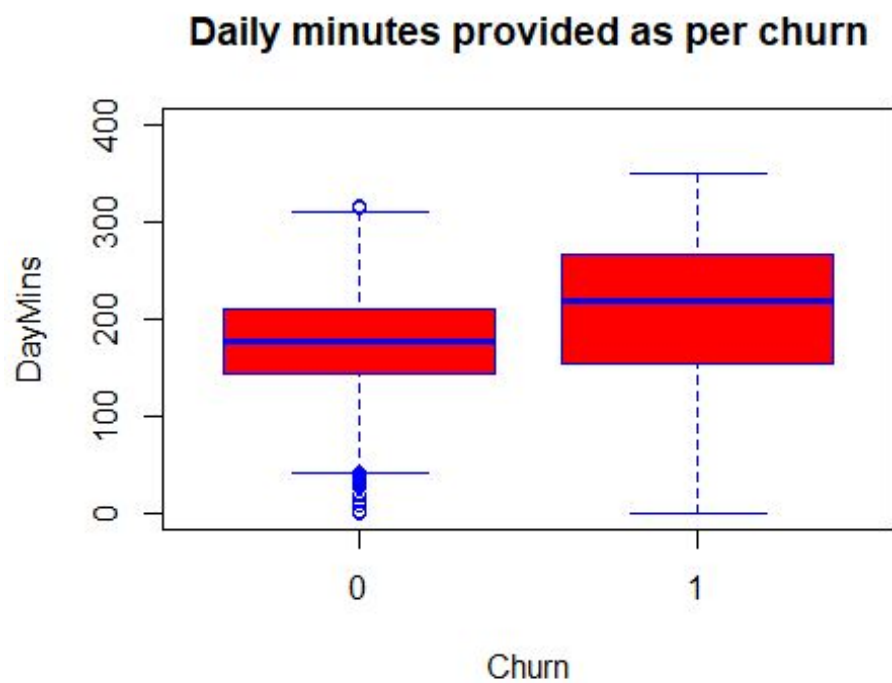
```



```

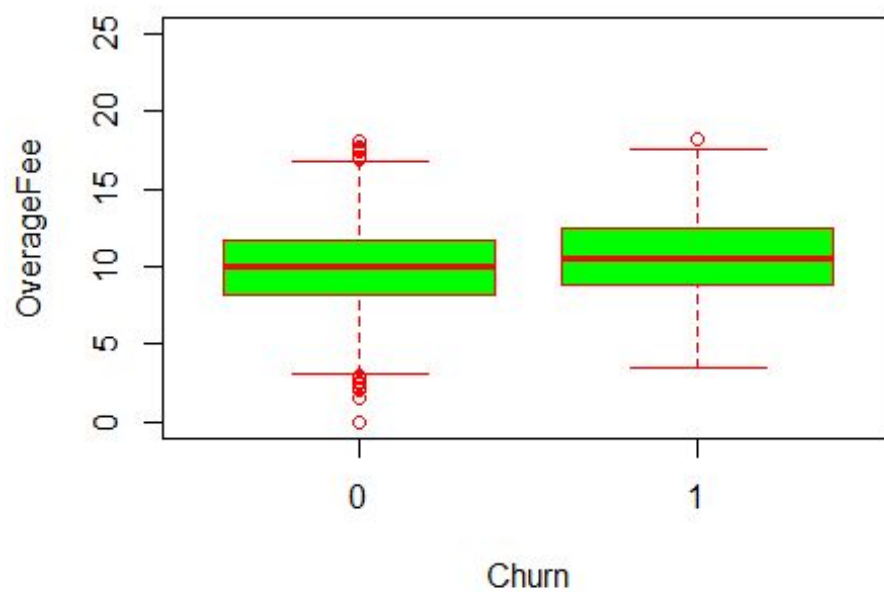
boxplot(`Min/Day`~Churn,
data = cell,
main = "Daily minutes provided as per churn",
xlab = "Churn",
ylab = "DayMins",
col = "Red",
border = "Blue",
ylim = c(0, 400)
)

```



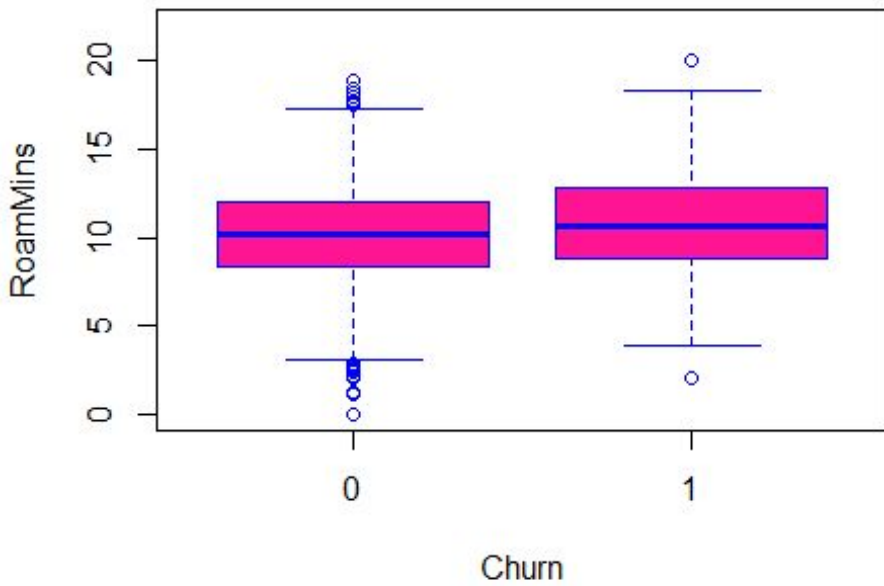
```
boxplot(Over.Fee~Churn,  
  data = cell,  
  main = "Churn based on the overage fee",  
  xlab = "Churn",  
  ylab = "OverageFee",  
  col = "Green",  
  border = "red",  
  ylim = c(0, 25)  
)
```

### Churn based on the overage fee




```
boxplot(RoamingMins~Churn,  
  data = cell,  
  main = "Churn Rate based on the roaming mins",  
  xlab = "Churn",  
  ylab = "RoamMins",  
  col = "Deeppink",  
  border = "blue",  
  ylim = c(0, 22)  
)
```


### Churn Rate based on the roaming mins



```
pairs.panels(cell[,c(2:11)],gap = 0, bg = c("red", "green","blue" , main =  
"Data Variations")[cell  
$Churn], pch = 21)
```




```
ggplot(cell, aes(y= AccWeeks, x = "", fill = Churn , col = "red")) +
  geom_boxplot()+
  theme_bw()+
  xlab(" ")
```



```
ggplot(cell, aes(y=`Min/Day`, x = "", fill = Churn )) +
  geom_boxplot()+
  theme_dark()+
  xlab(" ")
```





### **2.3.3. Insights gained from the Exploratory Data Analysis can be concluded as :-**

The variables in the dataset can be initially categorised into two types : categorical and Continuous ones.

#### **The categorical variables are :-**

- 1.Customer Service Calls
- 2.Contract Renewal
- 3.Data Plan

#### **The Continuous variables are :-**

- 1.AccountWeeks
- 2.Data Usage
- 3.Days/Min
- 4.Days/calls
- 5.Monthly Charges
- 6.Roaming Mins

There are other insights as well which are listed below that has been created from the EDA.

1.	Churn is the target variable for our analysis and we need to find the most optimized method of predicting the churn rate
2.	We do have outliers in the dataset in columns like : a)Churn, b)AccWeek, c)ContRenew , d)Usage , e)CScalls, f)Min/Day g)Call/Day h)Charge/Month i)over.fee j)RoamingMins.
3.	Multicollinearity does not exist in a huge amount amongst the variables. Most of them are either negatively correlated or has very minute correlation.

### 3. Model Creation

#### 3.1. Logistic Regression

Logistic regression is the appropriate regression analysis to conduct when the dependent variable is dichotomous (binary). Like all regression analyses, the logistic regression is a predictive analysis. Logistic regression is used to describe data and to explain the relationship between one dependent binary variable and one or more nominal, ordinal, interval or ratio-level independent variables.

*#We have selected churn as the target variable so detection of logistic regression be on that.*

```
cell$Churn = as.factor(cell$Churn)
cell$ContRenew = as.factor(cell$ContRenew)
cell$Plan = as.factor(cell$Plan)
one<- cell[which(cell$Churn == "1"),]
zero<- cell[which(cell$Churn == "0"),]
training1 <- sample(1:nrow(one),0.7*nrow(one))
training0 <- sample(1:nrow(zero),0.7*nrow(zero))
Final_training1 <- one[training1,]
Final_training0 <- zero[training0,]
trainingData <- rbind(Final_training1, Final_training0)
test1 <- one[-training1,]
test0 <- zero[-training0,]
testData <- rbind(test1, test0)
prop.table(table(testData$Churn))
```

```
##
##           0           1
## 0.8551449 0.1448551

Model1 <- glm(Churn ~ ., data = trainingData, family =
binomial(link="logit"))
summary(Model1)

##
## Call:
## glm(formula = Churn ~ ., family = binomial(link = "logit"), data =
trainingData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9959  -0.5080  -0.3419  -0.2015   3.0076
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -5.942045    0.661544  -8.982 < 2e-16 ***
## AccWeeks      -0.000276    0.001665  -0.166  0.86836
## ContRenew1    -1.994126    0.173088 -11.521 < 2e-16 ***
## Plan1         -1.097446    0.640829  -1.713  0.08680 .
## Usage         3.713872    2.315801   1.604  0.10878
## CScalls       0.494678    0.046809  10.568 < 2e-16 ***
## `Min/Day`     0.076251    0.039138   1.948  0.05138 .
## `Call/Day`    0.003849    0.003307   1.164  0.24451
## `Charge/Month` -0.365452    0.229898  -1.590  0.11192
## Over.Fee      0.759696    0.392729   1.934  0.05306 .
## RoamingMins   0.077428    0.026205   2.955  0.00313 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1930.1  on 2331  degrees of freedom
## Residual deviance: 1508.1  on 2321  degrees of freedom
## AIC: 1530.1
##
## Number of Fisher Scoring iterations: 6

anova(Model1, test = "Chisq")

## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Churn
##
## Terms added sequentially (first to last)
##
##
```

```
##           Df Deviance Resid. Df Resid. Dev Pr(>Chi)
## NULL                2331      1930.1
## AccWeeks           1    0.006      2330      1930.1 0.936623
## ContRenew           1  126.748      2329      1803.3 < 2.2e-16 ***
## Plan                1   31.150      2328      1772.2 2.389e-08 ***
## Usage              1    1.029      2327      1771.2 0.310382
## CScalls            1   98.258      2326      1672.9 < 2.2e-16 ***
## `Min/Day`          1  126.739      2325      1546.2 < 2.2e-16 ***
## `Call/Day`         1    1.218      2324      1545.0 0.269851
## `Charge/Month`     1   23.972      2323      1521.0 9.776e-07 ***
## Over.Fee           1    3.982      2322      1517.0 0.045998 *
## RoamingMins        1    8.905      2321      1508.1 0.002844 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
pR2(Model1)
```

```
##           llh      llhNull      G2      McFadden      r2ML
## -754.0494864 -965.0529427 422.0069125 0.2186444 0.1655342
##           r2CU
## 0.2940584
```

```
#checking the multicollinearity
car::vif(Model1)
```

```
##           AccWeeks      ContRenew      Plan      Usage
CScalls
##           1.006944      1.056386      13.858458      1598.838326
1.083021
##           `Min/Day`      `Call/Day` `Charge/Month`      Over.Fee
RoamingMins
##           915.488946      1.002070      2768.195614      214.073261
1.192887
```

```
#the model is affected by :Usage , `Min/Day` , `Charge/Month` , Over.Fee
```

```
#making the models by removing the above the list variables
```

```
mod1 = glm(Churn ~ . - Usage ,data = trainingData,family =
binomial(link="logit"))
summary(mod1)
```

```
##
## Call:
## glm(formula = Churn ~ . - Usage, family = binomial(link = "logit"),
##      data = trainingData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9995   -0.5030   -0.3412   -0.2035    3.0112
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -6.134660   0.651435  -9.417 < 2e-16 ***
```

```

## AccWeeks      -0.000144    0.001664   -0.087 0.931038
## ContRenew1    -1.987455    0.172900  -11.495 < 2e-16 ***
## Plan1         -0.973956    0.633689   -1.537 0.124303
## CScalls       0.494057    0.046695   10.581 < 2e-16 ***
## `Min/Day`     0.013820    0.003884    3.558 0.000374 ***
## `Call/Day`    0.003828    0.003300    1.160 0.246037
## `Charge/Month` 0.001591    0.021616    0.074 0.941328
## Over.Fee      0.134307    0.045424    2.957 0.003109 **
## RoamingMins   0.079861    0.026181    3.050 0.002286 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1930.1  on 2331  degrees of freedom
## Residual deviance: 1510.7  on 2322  degrees of freedom
## AIC: 1530.7
##
## Number of Fisher Scoring iterations: 6

anova(mod1,test = "Chisq")

## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Churn
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                      2331      1930.1
## AccWeeks          1      0.006      2330      1930.1  0.93662
## ContRenew          1    126.748      2329      1803.3 < 2.2e-16 ***
## Plan              1     31.150      2328      1772.2 2.389e-08 ***
## CScalls            1     97.703      2327      1674.5 < 2.2e-16 ***
## `Min/Day`          1    126.469      2326      1548.0 < 2.2e-16 ***
## `Call/Day`          1     1.200      2325      1546.8  0.27323
## `Charge/Month`      1     22.439      2324      1524.4 2.169e-06 ***
## Over.Fee           1      4.207      2323      1520.2  0.04025 *
## RoamingMins         1      9.505      2322      1510.7  0.00205 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

mod2 = glm(Churn ~ . - `Min/Day`, data = trainingData, family =
binomial(link="logit"))
summary(mod2)

##
## Call:
## glm(formula = Churn ~ . - `Min/Day`, family = binomial(link = "logit"),

```

```
##      data = trainingData)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q        Max
## -1.9967   -0.5029   -0.3413   -0.2028    3.0058
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -6.1470698   0.6538640  -9.401  < 2e-16 ***
## AccWeeks      -0.0001181   0.0016630  -0.071  0.943377
## ContRenew1    -1.9845677   0.1727693 -11.487  < 2e-16 ***
## Plan1         -1.0593968   0.6392541  -1.657  0.097471 .
## Usage         -0.7775758   0.2298553  -3.383  0.000717 ***
## CScalls        0.4941528   0.0466706   10.588  < 2e-16 ***
## `Call/Day`     0.0038332   0.0032973    1.163  0.245019
## `Charge/Month` 0.0823911   0.0077398   10.645  < 2e-16 ***
## Over.Fee      -0.0033558   0.0293955  -0.114  0.909111
## RoamingMins    0.0784603   0.0261994    2.995  0.002747 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1930.1  on 2331  degrees of freedom
## Residual deviance: 1511.9  on 2322  degrees of freedom
## AIC: 1531.9
##
## Number of Fisher Scoring iterations: 6
```

```
anova(mod2,test = "Chisq")
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Churn
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                                2331      1930.1
## AccWeeks          1      0.006      2330      1930.1  0.936623
## ContRenew         1    126.748      2329      1803.3 < 2.2e-16 ***
## Plan              1     31.150      2328      1772.2 2.389e-08 ***
## Usage             1      1.029      2327      1771.2  0.310382
## CScalls           1     98.258      2326      1672.9 < 2.2e-16 ***
## `Call/Day`        1      1.406      2325      1671.5  0.235760
## `Charge/Month`    1    150.397      2324      1521.1 < 2.2e-16 ***
## Over.Fee          1      0.042      2323      1521.1  0.838191
## RoamingMins       1      9.158      2322      1511.9  0.002476 **
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

mod3 = glm(Churn ~ . - `Charge/Month`, data = trainingData, family =
binomial(link="logit"))
summary(mod3)

##
## Call:
## glm(formula = Churn ~ . - `Charge/Month`, family = binomial(link =
"logit"),
##     data = trainingData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9975  -0.5035  -0.3408  -0.2028   3.0102
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -6.1234718  0.6524488  -9.385  < 2e-16 ***
## AccWeeks    -0.0001468  0.0016632  -0.088  0.92965
## ContRenew1  -1.9862388  0.1728542 -11.491  < 2e-16 ***
## Plan1       -1.0671811  0.6396459  -1.668  0.09524 .
## Usage        0.0489381  0.2177992   0.225  0.82222
## CScalls      0.4945112  0.0467088  10.587  < 2e-16 ***
## `Min/Day`    0.0140941  0.0013181  10.692  < 2e-16 ***
## `Call/Day`   0.0038354  0.0032997   1.162  0.24508
## Over.Fee     0.1370751  0.0271191   5.055 4.31e-07 ***
## RoamingMins  0.0782991  0.0262027   2.988  0.00281 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1930.1  on 2331  degrees of freedom
## Residual deviance: 1510.6  on 2322  degrees of freedom
## AIC: 1530.6
##
## Number of Fisher Scoring iterations: 6

anova(mod3, test = "Chisq")

## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Churn
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
```

```
## NULL                2331      1930.1
## AccWeeks           1      0.006      2330      1930.1  0.936623
## ContRenew          1  126.748      2329      1803.3 < 2.2e-16 ***
## Plan                1   31.150      2328      1772.2 2.389e-08 ***
## Usage               1    1.029      2327      1771.2  0.310382
## CScalls             1   98.258      2326      1672.9 < 2.2e-16 ***
## `Min/Day`          1  126.739      2325      1546.2 < 2.2e-16 ***
## `Call/Day`         1    1.218      2324      1545.0  0.269851
## Over.Fee           1   25.207      2323      1519.8 5.149e-07 ***
## RoamingMins        1    9.116      2322      1510.6  0.002533 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

mod4 = glm(Churn ~ . - Over.Fee , data = trainingData, family =
binomial(link="logit"))
summary(mod4)

##
## Call:
## glm(formula = Churn ~ . - Over.Fee, family = binomial(link = "logit"),
##      data = trainingData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9954  -0.5037  -0.3410  -0.2031   3.0075
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -6.1211380   0.6555746  -9.337  < 2e-16 ***
## AccWeeks      -0.0001282   0.0016631  -0.077  0.93858
## ContRenew1    -1.9838356   0.1727524 -11.484  < 2e-16 ***
## Plan1         -1.0582243   0.6392634  -1.655  0.09785 .
## Usage         -0.7372190   0.2677779  -2.753  0.00590 **
## CScalls        0.4940484   0.0466715  10.586  < 2e-16 ***
## `Min/Day`      0.0007632   0.0029296   0.261  0.79447
## `Call/Day`     0.0038231   0.0032976   1.159  0.24631
## `Charge/Month` 0.0783494   0.0158639   4.939 7.86e-07 ***
## RoamingMins    0.0783426   0.0262022   2.990  0.00279 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1930.1  on 2331  degrees of freedom
## Residual deviance: 1511.9  on 2322  degrees of freedom
## AIC: 1531.9
##
## Number of Fisher Scoring iterations: 6

anova(mod4, test = "Chisq")
```



```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Churn
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                      2331      1930.1
## AccWeeks          1    0.006      2330      1930.1  0.936623
## ContRenew          1  126.748      2329      1803.3 < 2.2e-16 ***
## Plan              1   31.150      2328      1772.2 2.389e-08 ***
## Usage             1    1.029      2327      1771.2  0.310382
## CScalls           1   98.258      2326      1672.9 < 2.2e-16 ***
## `Min/Day`         1  126.739      2325      1546.2 < 2.2e-16 ***
## `Call/Day`        1    1.218      2324      1545.0  0.269851
## `Charge/Month`    1   23.972      2323      1521.0 9.776e-07 ***
## RoamingMins       1    9.129      2322      1511.9  0.002516 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#predicting the models on the training dataset
pred<-predict(Model1,newdata=trainingData,type="response")
prediction<- ifelse(pred>0.5,1,0)
prediction1 <- factor(prediction, levels=c(0,1))
act <- trainingData$Churn
confusionMatrix(prediction1,act,positive="1")
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 1942  269
##              1   52   69
##
##              Accuracy : 0.8623
##              95% CI : (0.8477, 0.8761)
##              No Information Rate : 0.8551
##              P-Value [Acc > NIR] : 0.166
##
##              Kappa : 0.2428
##
##              Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.20414
##              Specificity : 0.97392
##              Pos Pred Value : 0.57025
##              Neg Pred Value : 0.87834
##              Prevalence : 0.14494
##              Detection Rate : 0.02959
```

```

##      Detection Prevalence : 0.05189
##      Balanced Accuracy : 0.58903
##
##      'Positive' Class : 1
##

pred1<-predict(mod1,newdata=trainingData,type="response")
prediction1<- ifelse(pred1>0.5,1,0)
prediction2 <- factor(prediction1, levels=c(0,1))
act <- trainingData$Churn
confusionMatrix(prediction2,act,positive="1")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##      0 1946  265
##      1   48   73
##
##              Accuracy : 0.8658
##              95% CI : (0.8513, 0.8794)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.07367
##
##              Kappa : 0.2617
##
##  Mcnemar's Test P-Value : < 2e-16
##
##              Sensitivity : 0.21598
##              Specificity : 0.97593
##              Pos Pred Value : 0.60331
##              Neg Pred Value : 0.88014
##              Prevalence : 0.14494
##              Detection Rate : 0.03130
##      Detection Prevalence : 0.05189
##      Balanced Accuracy : 0.59595
##
##      'Positive' Class : 1
##

pred2<-predict(mod2,newdata=trainingData,type="response")
prediction3<- ifelse(pred2>0.5,1,0)
prediction4 <- factor(prediction3, levels=c(0,1))
act <- trainingData$Churn
confusionMatrix(prediction4,act,positive="1")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##      0 1945  268
##      1   49   70

```

```

##
##          Accuracy : 0.8641
##          95% CI : (0.8495, 0.8777)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.1133
##
##          Kappa : 0.2497
##
##  McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.20710
##          Specificity : 0.97543
##      Pos Pred Value : 0.58824
##      Neg Pred Value : 0.87890
##          Prevalence : 0.14494
##      Detection Rate : 0.03002
##      Detection Prevalence : 0.05103
##      Balanced Accuracy : 0.59126
##
##      'Positive' Class : 1
##

pred3<-predict(mod3,newdata=trainingData,type="response")
prediction5<- ifelse(pred3>0.5,1,0)
prediction6 <- factor(prediction5, levels=c(0,1))
act <- trainingData$Churn
confusionMatrix(prediction6,act,positive="1")

## Confusion Matrix and Statistics
##
##          Reference
## Prediction    0    1
##          0 1947  266
##          1   47   72
##
##          Accuracy : 0.8658
##          95% CI : (0.8513, 0.8794)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.07367
##
##          Kappa : 0.2592
##
##  McNemar's Test P-Value : < 2e-16
##
##          Sensitivity : 0.21302
##          Specificity : 0.97643
##      Pos Pred Value : 0.60504
##      Neg Pred Value : 0.87980
##          Prevalence : 0.14494
##      Detection Rate : 0.03087
##      Detection Prevalence : 0.05103
##      Balanced Accuracy : 0.59472

```

```
##
##      'Positive' Class : 1
##

pred4<-predict(mod4,newdata=trainingData,type="response")
prediction7<- ifelse(pred4>0.5,1,0)
prediction8 <- factor(prediction7, levels=c(0,1))
act <- trainingData$Churn
confusionMatrix(prediction8,act,positive="1")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##           0 1946  268
##           1   48   70
##
##              Accuracy : 0.8645
##              95% CI : (0.8499, 0.8781)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.1022
##
##              Kappa : 0.2508
##
##  Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.20710
##              Specificity : 0.97593
##              Pos Pred Value : 0.59322
##              Neg Pred Value : 0.87895
##              Prevalence : 0.14494
##              Detection Rate : 0.03002
##      Detection Prevalence : 0.05060
##              Balanced Accuracy : 0.59151
##
##      'Positive' Class : 1
##

#After training data is tested we can go forward and start with the test data
pred_test<-predict(Model1,newdata=testData,type="response")
prediction_test<- ifelse(pred_test>0.5,1,0)
prediction1_test <- factor(prediction_test, levels=c(0,1))
act <- testData$Churn
confusionMatrix(prediction1_test,act,positive="1")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##           0  837 127
##           1   19  18
```

```

##
##          Accuracy : 0.8541
##          95% CI : (0.8307, 0.8754)
##    No Information Rate : 0.8551
##    P-Value [Acc > NIR] : 0.5577
##
##          Kappa : 0.1476
##
##  McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.12414
##          Specificity : 0.97780
##          Pos Pred Value : 0.48649
##          Neg Pred Value : 0.86826
##          Prevalence : 0.14486
##          Detection Rate : 0.01798
##          Detection Prevalence : 0.03696
##          Balanced Accuracy : 0.55097
##
##          'Positive' Class : 1
##

pred1_test<-predict(mod1,newdata=testData,type="response")
prediction1_test<- ifelse(pred1_test>0.5,1,0)
prediction2_test <- factor(prediction1_test, levels=c(0,1))
act <- testData$Churn
confusionMatrix(prediction2_test,act,positive="1")

## Confusion Matrix and Statistics
##
##          Reference
## Prediction    0    1
##          0 836 126
##          1  20  19
##
##          Accuracy : 0.8541
##          95% CI : (0.8307, 0.8754)
##    No Information Rate : 0.8551
##    P-Value [Acc > NIR] : 0.5577
##
##          Kappa : 0.1546
##
##  McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.13103
##          Specificity : 0.97664
##          Pos Pred Value : 0.48718
##          Neg Pred Value : 0.86902
##          Prevalence : 0.14486
##          Detection Rate : 0.01898
##          Detection Prevalence : 0.03896
##          Balanced Accuracy : 0.55383

```

```
##
##      'Positive' Class : 1
##

pred2_test<-predict(mod2,newdata=testData,type="response")
prediction3_test<- ifelse(pred2_test>0.5,1,0)
prediction4_test <- factor(prediction3_test, levels=c(0,1))
act <- testData$Churn
confusionMatrix(prediction4_test,act,positive="1")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 837 126
##              1   19   19
##
##              Accuracy : 0.8551
##              95% CI : (0.8318, 0.8764)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.5221
##
##              Kappa : 0.1569
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.13103
##              Specificity : 0.97780
##              Pos Pred Value : 0.50000
##              Neg Pred Value : 0.86916
##              Prevalence : 0.14486
##              Detection Rate : 0.01898
##      Detection Prevalence : 0.03796
##              Balanced Accuracy : 0.55442
##
##      'Positive' Class : 1
##

pred3_test<-predict(mod3,newdata=testData,type="response")
prediction5_test<- ifelse(pred3_test>0.5,1,0)
prediction6_test <- factor(prediction5_test, levels=c(0,1))
act <- testData$Churn
confusionMatrix(prediction6_test,act,positive="1")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 836 126
##              1   20   19
##
##              Accuracy : 0.8541
```

```

##          95% CI : (0.8307, 0.8754)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.5577
##
##          Kappa : 0.1546
##
##  McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.13103
##          Specificity : 0.97664
##          Pos Pred Value : 0.48718
##          Neg Pred Value : 0.86902
##          Prevalence : 0.14486
##          Detection Rate : 0.01898
##          Detection Prevalence : 0.03896
##          Balanced Accuracy : 0.55383
##
##          'Positive' Class : 1
##


pred4_test<-predict(mod4,newdata=testData,type="response")
prediction7_test<- ifelse(pred4_test>0.5,1,0)
prediction8_test <- factor(prediction7_test, levels=c(0,1))
act <- testData$Churn
confusionMatrix(prediction8_test,act,positive="1")

## Confusion Matrix and Statistics
##
##          Reference
## Prediction    0    1
##          0 836 126
##          1  20  19
##
##          Accuracy : 0.8541
##          95% CI : (0.8307, 0.8754)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.5577
##
##          Kappa : 0.1546
##
##  McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.13103
##          Specificity : 0.97664
##          Pos Pred Value : 0.48718
##          Neg Pred Value : 0.86902
##          Prevalence : 0.14486
##          Detection Rate : 0.01898
##          Detection Prevalence : 0.03896
##          Balanced Accuracy : 0.55383
##

```

```
##          'Positive' Class : 1
##

#ROC plot and AUC of the different models we have created in Logistic Regression
rocpred<-predict(Model1,testData,type = 'response')
rocpred<-prediction(pred_test, testData$Churn)
roc<-performance(rocpred,"tpr", "fpr")
plot(roc)
```




```
auc<-performance(rocpred,"auc")
auc

## An object of class "performance"
## Slot "x.name":
## [1] "None"
##
## Slot "y.name":
## [1] "Area under the ROC curve"
##
## Slot "alpha.name":
## [1] "none"
##
## Slot "x.values":
## list()
##
## Slot "y.values":
## [[1]]
## [1] 0.7995488
```



```
##
##
## Slot "alpha.values":
## list()

rocpred_mod1<-predict(mod1,testData,type = 'response')
rocpred_mod1<-prediction(pred1_test, testData$Churn)
roc_mod1<-performance(rocpred_mod1,"tpr", "fpr")
plot(roc_mod1)
```




```
auc_mod1<-performance(rocpred_mod1,"auc")
auc_mod1

## An object of class "performance"
## Slot "x.name":
## [1] "None"
##
## Slot "y.name":
## [1] "Area under the ROC curve"
##
## Slot "alpha.name":
## [1] "none"
##
## Slot "x.values":
## list()
##
## Slot "y.values":
## [[1]]
## [1] 0.8038269
```

```
##
##
## Slot "alpha.values":
## list()

rocpred_mod2<-predict(mod2,testData,type = 'response')
rocpred_mod2<-prediction(pred2_test, testData$Churn)
roc_mod2<-performance(rocpred_mod2,"tpr", "fpr")
plot(roc_mod2)
```




```
auc_mod2<-performance(rocpred_mod2,"auc")
auc_mod2
```

```
## An object of class "performance"
## Slot "x.name":
## [1] "None"
##
## Slot "y.name":
## [1] "Area under the ROC curve"
##
## Slot "alpha.name":
## [1] "none"
##
## Slot "x.values":
## list()
##
## Slot "y.values":
## [[1]]
## [1] 0.8048501
```

```
##
##
## Slot "alpha.values":
## list()

rocpred_mod3<-predict(mod3,testData,type = 'response')
rocpred_mod3<-prediction(pred3_test, testData$Churn)
roc_mod3<-performance(rocpred_mod3,"tpr", "fpr")
plot(roc_mod3)
```




```
auc_mod3<-performance(rocpred_mod3,"auc")
auc_mod3

## An object of class "performance"
## Slot "x.name":
## [1] "None"
##
## Slot "y.name":
## [1] "Area under the ROC curve"
##
## Slot "alpha.name":
## [1] "none"
##
## Slot "x.values":
## list()
##
## Slot "y.values":
## [[1]]
## [1] 0.803972
```

```
##
##
## Slot "alpha.values":
## list()

rocpred_mod4<-predict(mod4,testData,type = 'response')
rocpred_mod4<-prediction(pred4_test, testData$Churn)
roc_mod4<-performance(rocpred_mod4,"tpr", "fpr")
plot(roc_mod4)
```



```
auc_mod4<-performance(rocpred_mod4,"auc")
auc_mod4

## An object of class "performance"
## Slot "x.name":
## [1] "None"
##
## Slot "y.name":
## [1] "Area under the ROC curve"
##
## Slot "alpha.name":
## [1] "none"
##
## Slot "x.values":
## list()
##
## Slot "y.values":
## [[1]]
## [1] 0.8046407
```

```
##
##
## Slot "alpha.values":
## list()
```

So we can see that almost all the graphs are nearly the same for these models. The model1 is having all the elements while the other 4 models doesnot have the elements which were having high collinearity.  
#The next model will not be having the highly collinear elements and we would examine the model.

```
Model2 = glm(Churn ~ AccWeeks + ContRenew + CScalls + `Call/Day` +
RoamingMins ,data = trainingData , family = binomial(link = logit))
summary(Model2)
```

```
##
## Call:
## glm(formula = Churn ~ AccWeeks + ContRenew + CScalls + `Call/Day` +
##      RoamingMins, family = binomial(link = logit), data = trainingData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5896  -0.5225  -0.4228  -0.3414   2.5721
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.9338635   0.4542735  -4.257 2.07e-05 ***
## AccWeeks    -0.0004651   0.0015919  -0.292  0.77015
## ContRenew1  -1.9234541   0.1593943 -12.067 < 2e-16 ***
## CScalls      0.4277921   0.0429859   9.952 < 2e-16 ***
## `Call/Day`   0.0035614   0.0031207   1.141  0.25378
## RoamingMins  0.0667935   0.0229942   2.905  0.00367 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1930.1  on 2331  degrees of freedom
## Residual deviance: 1694.2  on 2326  degrees of freedom
## AIC: 1706.2
##
## Number of Fisher Scoring iterations: 5

anova(Model2,test = "Chisq")

## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Churn
##
## Terms added sequentially (first to last)
```

```
##
##
##           Df Deviance Resid. Df Resid. Dev Pr(>Chi)
## NULL                2331      1930.1
## AccWeeks      1      0.006      2330      1930.1 0.93662
## ContRenew     1    126.748      2329      1803.3 < 2e-16 ***
## CScalls       1     99.065      2328      1704.3 < 2e-16 ***
## `Call/Day`    1      1.471      2327      1702.8 0.22513
## RoamingMins   1      8.568      2326      1694.2 0.00342 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

*#Checking the multi collinearity*

```
car::vif(Model2)
```

```
##      AccWeeks      ContRenew      CScalls      `Call/Day`      RoamingMins
##      1.001538      1.033876      1.034264      1.001396      1.002103
```

*#Prediction using the model for the training dataset*

```
pred_Model2<-predict(Model2,newdata=trainingData,type="response")
prediction_Model2<- ifelse(pred_Model2>0.5,1,0)
prediction1_Model2 <- factor(prediction_Model2, levels=c(0,1))
act <- trainingData$Churn
confusionMatrix(prediction1_Model2,act,positive="1")
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction    0    1
##           0 1962  295
##           1   32   43
##
##           Accuracy : 0.8598
##           95% CI : (0.845, 0.8736)
##           No Information Rate : 0.8551
##           P-Value [Acc > NIR] : 0.2698
##
##           Kappa : 0.1642
##
##  Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.12722
##           Specificity : 0.98395
##           Pos Pred Value : 0.57333
##           Neg Pred Value : 0.86930
##           Prevalence : 0.14494
##           Detection Rate : 0.01844
##           Detection Prevalence : 0.03216
##           Balanced Accuracy : 0.55559
##
##           'Positive' Class : 1
##
```


```

#Prediction using the test dataset
pred_Model2_test<-predict(Model2,newdata=testData,type="response")
prediction_Model2_test<- ifelse(pred_Model2_test>0.5,1,0)
prediction1_Model2_test <- factor(prediction_Model2_test, levels=c(0,1))
act <- testData$Churn
confusionMatrix(prediction1_Model2_test,act,positive="1")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 843 132
##              1  13  13
##
##              Accuracy : 0.8551
##              95% CI : (0.8318, 0.8764)
##      No Information Rate : 0.8551
##      P-Value [Acc > NIR] : 0.5221
##
##              Kappa : 0.113
##
##  Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.08966
##              Specificity : 0.98481
##              Pos Pred Value : 0.50000
##              Neg Pred Value : 0.86462
##              Prevalence : 0.14486
##              Detection Rate : 0.01299
##      Detection Prevalence : 0.02597
##      Balanced Accuracy : 0.53723
##
##              'Positive' Class : 1
##

#ROC Curve and AUC
rocpred_Model2<-predict(Model2,testData,type = 'response')
rocpred_Model2<-prediction(pred_Model2_test, testData$Churn)
roc_Model2<-performance(rocpred_Model2,"tpr", "fpr")
plot(roc_Model2)

```

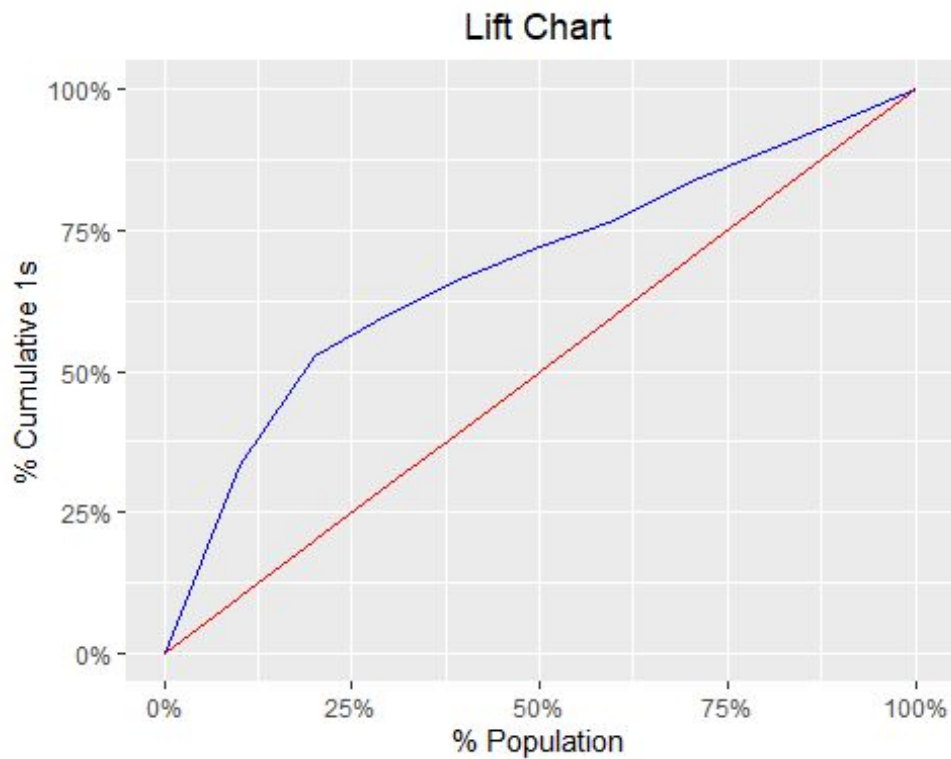


```
auc<-performance(rocpred_Model12,"auc")
auc
```

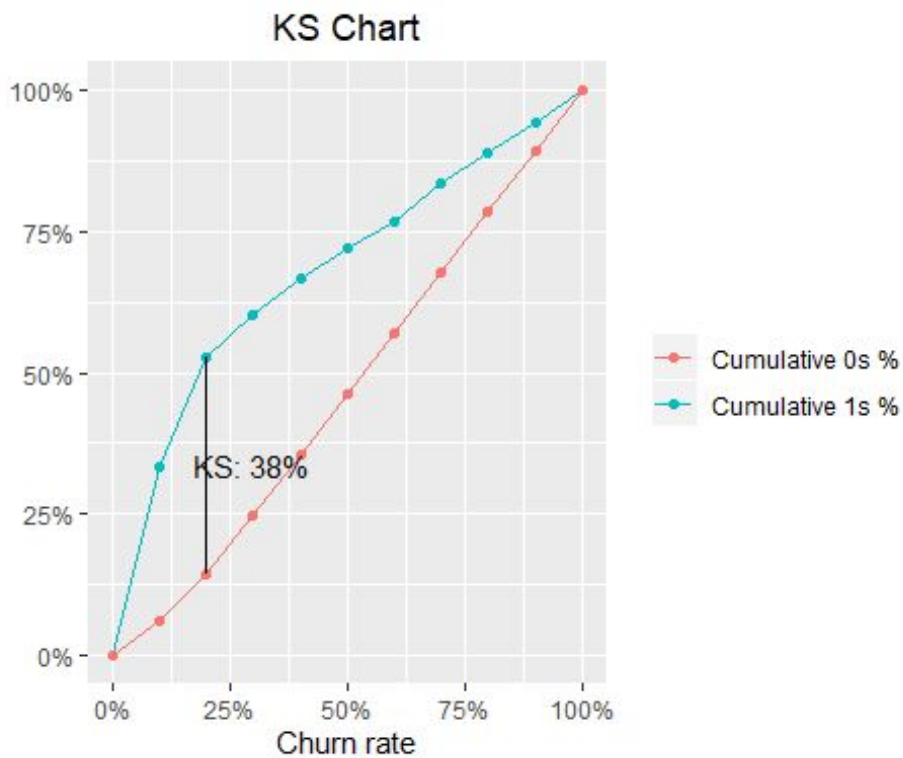
```
## An object of class "performance"
## Slot "x.name":
## [1] "None"
##
## Slot "y.name":
## [1] "Area under the ROC curve"
##
## Slot "alpha.name":
## [1] "none"
##
## Slot "x.values":
## list()
##
## Slot "y.values":
## [[1]]
## [1] 0.7214953
##
##
## Slot "alpha.values":
## list()
```

```
k = blr_gains_table(Model12)
plot(k)
```

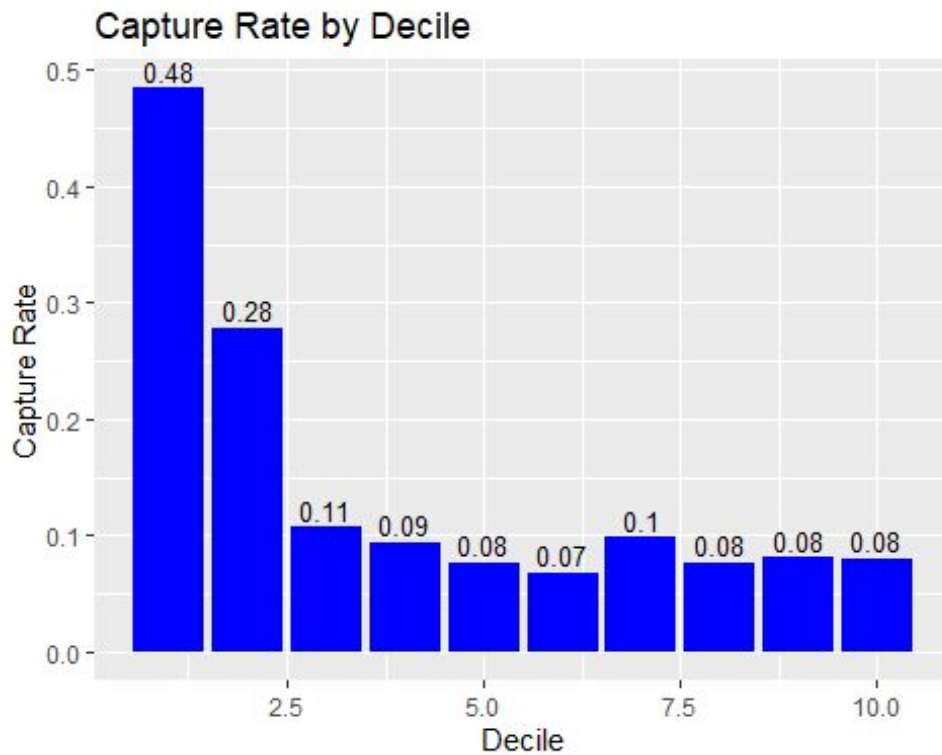





```
blr_ks_chart(k, title = "KS Chart",  
            yaxis_title = " ", xaxis_title = "Churn rate",  
            ks_line_color = "black")
```




```
blr_decile_lift_chart(k, xaxis_title = "Decile",
                      yaxis_title = "Decile Mean / Global Mean",
                      title = "Decile Lift Chart",
                      bar_color = "blue", text_size = 3.5,
                      text_vjust = -0.3)
```



```
blr_roc_curve(k, title = "ROC Curve",
              xaxis_title = "1 - Specificity",
              yaxis_title = "Sensitivity", roc_curve_col = "blue",
              diag_line_col = "red", point_shape = 18,
              point_fill = "blue", point_color = "blue",
              plot_title_justify = 0.5)
```



```
blr_plot_difchisq_fitted(Model12, point_color = "blue",  
                           title = "Delta Chi Square vs Fitted Values Plot",  
                           xaxis_title = "Fitted Values",  
                           yaxis_title = "Delta Chi Square")
```



### 3.2.K-Nearest Neighbour(KNN) :-

The k-nearest neighbors (**KNN**) algorithm is a simple, easy-to-implement supervised **machine learning** algorithm that can be used to solve both classification and regression problems.

```
cell$ContRenew = as.numeric(cell$ContRenew)
cell$Plan = as.numeric(cell$Plan)

#Normalizing the data
normalize<-function(x){ return((x-min(x))/(max(x)-min(x)))
}

cell.norm<-as.data.frame(lapply(cell[,-1],normalize ))
#View(cell.norm)
usable.data = cbind(cell[,1], cell.norm)
str(usable.data)

## 'data.frame':    3333 obs. of  11 variables:
## $ Churn          : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ AccWeeks       : num  0.525 0.438 0.562 0.343 0.306 ...
## $ ContRenew      : num  1 1 1 0 0 0 1 0 1 0 ...
## $ Plan           : num  1 1 0 0 0 0 1 0 0 1 ...
## $ Usage          : num  0.5 0.685 0 0 0 ...
## $ CScalls        : num  0.111 0.111 0 0.222 0.333 ...
## $ Min.Day        : num  0.756 0.461 0.694 0.853 0.475 ...
## $ Call.Day       : num  0.667 0.745 0.691 0.43 0.685 ...
## $ Charge.Month   : num  0.771 0.699 0.391 0.442 0.277 ...
## $ Over.Fee       : num  0.543 0.538 0.333 0.17 0.408 ...
## $ RoamingMins    : num  0.5 0.685 0.61 0.33 0.505 0.315 0.375 0.355 0.435
0.56 ...

#View(usable.data)
# Data partitioning
spl = sample.split(usable.data$Churn, SplitRatio = 0.7)
train = subset(usable.data, spl == T)
test = subset(usable.data, spl == F)
dim(train)

## [1] 2333    11

dim(test)

## [1] 1000    11

pred_knn_5 = knn(train[-1], test[-1], train[,1], k = 5)
table.knn_5 = table(test[,1], pred_knn_5)
accuracy_knn_5 = sum(diag(table.knn_5)/sum(table.knn_5))
accuracy_knn_5

## [1] 0.904
```

```

loss.knn.5<-table.knn_5[2,1]/(table.knn_5[2,1]+table.knn_5[1,1])
loss.knn.5

## [1] 0.09061489

pred_knn_9 = knn(train[-1], test[-1], train[,1], k = 9)
table.knn_9 = table(test[,1], pred_knn_9)
accuracy_knn_9 = sum(diag(table.knn_9)/sum(table.knn_9))
accuracy_knn_9

## [1] 0.898

loss.knn.9<-table.knn_9[2,1]/(table.knn_9[2,1]+table.knn_9[1,1])
loss.knn.9

## [1] 0.0981857

```

### 3.3.Naive Bayes


Naive Bayes classifiers are a collection of classification algorithms based on **Bayes' Theorem**. It is not a single algorithm but a family of algorithms where all of them share a common principle, i.e. every pair of features being classified is independent of each other.

```

train.2fact = train[,c(7,8,9)]
val.2fact = train[,c(7,8,9)]
set = train.2fact
X1 = seq(min(set[, 1])-0.1 , max(set[, 1])+0.1 , by = 0.005)
X2 = seq(min(set[, 2]) -0.1, max(set[, 2])+0.1 , by = 0.005)
grid_set = expand.grid(X1, X2)
colnames(grid_set) = c('mean.assets', 'mean.studied')
plot(set[, -3],
      main = 'Naive Bayes (Training set)',
      xlab = 'mean.assets', ylab = 'mean.studied',
      xlim = range(X1), ylim = range(X2))


```

## Naive Bayes (Training set)



```
set = val.2fact
X1 = seq(min(set[, 1])-0.1 , max(set[, 1])+0.1 , by = 0.005)
X2 = seq(min(set[, 2]) -0.1, max(set[, 2])+0.1 , by = 0.005)
grid_set = expand.grid(X1, X2)
colnames(grid_set) = c('mean.assets', 'mean.studied')
plot(set[, -3],
      main = 'Naive Bayes (Val set)',
      xlab = 'meanassets', ylab = 'mean.studied',
      xlim = range(X1), ylim = range(X2))
```

## Naive Bayes (Val set)



```
NB = naiveBayes(Churn ~., data = train)
predNB = predict(NB, test, type = "class")
tab.NB = table(test[,1], predNB)
tab.NB

##      predNB
##         0   1
##  0 801  54
##  1   84  61

accuracy_NB = sum(diag(tab.NB)/sum(tab.NB))
accuracy_NB

## [1] 0.862

#Confusion Matrix
confusion.matrix(predNB, test$Churn)

## Warning in Ops.factor(pred, threshold): '>=' not meaningful for factors
## Warning in Ops.factor(pred, threshold): '<' not meaningful for factors

##      obs
## pred  0   1
##   0 801  54
##   1   84  61
## attr(,"class")
## [1] "confusion.matrix"
```

## 4. Model Performance

Model Evaluation is an integral part of the model development process. It helps to find the best model that represents our data and how well the chosen model will work in the future.

Evaluating model performance with the data used for training is not acceptable in data science because it can easily generate overoptimistic and overfitted models.

There are two methods of evaluating models in data science, Hold-Out and Cross-Validation. To avoid overfitting, both methods use a test set (not seen by the model) to evaluate model performance.

```
splitSample <- sample(1:2, size = nrow(cell), prob = c(0.7,
0.3), replace = T)

train_set <- cell[splitSample == 1, ]

intrain <- sample(1:2, size = nrow(train_set), prob =
c(0.7, 0.3), replace = T)

trainset <- train_set[intrain == 1, ]


validset <- train_set[intrain == 2, ]

testset <- cell[splitSample == 2, ]

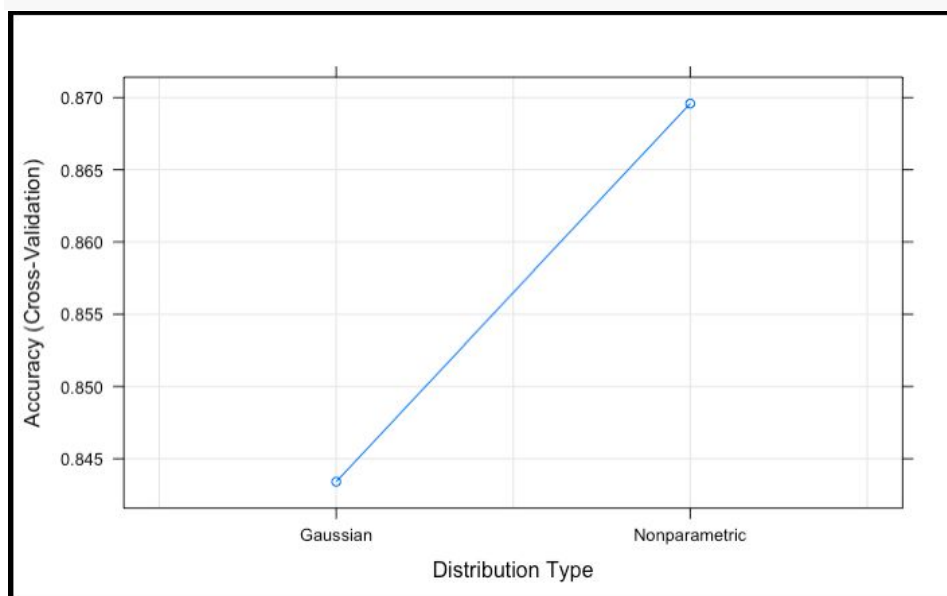
#cross validation of the data
tcontrol <- trainControl(method = "cv", number = 10)
set.seed(1234)

# KNN
modelKNN <- train(Churn ~ ., data = trainset, method =
"knn", preProcess = c("center",
"scale"), trControl = tcontrol) # data is normalised
using Preprocess
par(mfrow = c(1,2))
plot(modelKNN)
```









```
modelNB <- train(Churn ~ ., data = trainset, method = "nb",
trControl = tcontrol)
```



```
modelLG <- train(Churn ~ ., data = trainset, method = "glm", family = binomial,
trControl = tcontrol)
plot(modelLG$finalModel)
```





```
# KNN
```

```
pKNN = predict(modelKNN, validset)
pKNN_test = predict(modelKNN, testset)
```

```
# Naive Bayes
```

```
pNB = predict(modelNB, validset)
pNB_test = predict(modelNB, testset)
```

```
# Logistic Regression
```

```
pLG = predict(modelLG, validset)
pLG_test = predict(modelLG, testset)
```

```
# KNN
```

```
cmKNN = confusionMatrix(validset$Churn, pKNN)
cmKNN_test = confusionMatrix(testset$Churn, pKNN_test)
```

```
# Naive Bayes
```

```
cmNB = confusionMatrix(validset$Churn, pNB)
cmNB_test = confusionMatrix(testset$Churn, pNB_test)
```

```
# Logistic Regression
```

```
cmLG <- confusionMatrix(validset$Churn, pLG)
cmLG = confusionMatrix(testset$Churn, pLG_test)
```

```
ModelType <- c("K nearest neighbor", "Naive Bayes",
               "Logistic regression")
```

```
# classification accuracy
```

```
TrainAccuracy <- c(max(modelKNN$results$Accuracy),
                  max(modelNB$results$Accuracy),
                  max(modelLG$results$Accuracy))
```

```
# Training misclassification error
```

```
Train_missclass_Error <- 1 - TrainAccuracy
```

```
# validation classification accuracy
```

```
ValidationAccuracy <- c(cmKNN$overall[1], cmNB$overall[1],
                       cmLG$overall[1])
```

```
# Validation misclassification error or out-of-sample-error
Validation_missclass_Error <- 1 - ValidationAccuracy

metrics <- data.frame(ModelType, TrainAccuracy,
  Train_missclass_Error, ValidationAccuracy,
  Validation_missclass_Error) # data frame with above
metrics

knitr::kable(metrics, digits = 5) # print table using
kable() from knitr package
```

ModelType	TrainAccuracy	Train_missclass_Error	ValidationAccuracy	Validation_missclass_Error
K nearest neighbor	0.90089	0.09911	0.89787	0.10213
Naive Bayes	0.87312	0.12688	0.88227	0.11773
Logistic regression	0.85621	0.14379	0.86845	0.13155

Accuracy and Sensitivity is relatively high for KNN among the above methods. Yet, insights from logistic regression model can still be utilized to assist decision makers.

An organization loses its customers to its competition for various reasons. Churn can affect the company's overall growth. The list of the factors that affect the most are listed as :

- 1.Contract Renewal
2. Data Plan
- 3.Customer Service Calls
- 4.Day Mins
- 5.Overage Fee
- 6.Roaming Minutes

Keeping a strict check on the above factors we can reduce the churn rate amongst the customers.

